

COMP 530

Introduction to Operating Systems

Notes on HW 4

Implementing Producer/Consumer systems With Message Passing

Kevin Jeffay
Department of Computer Science
University of North Carolina at Chapel Hill
jeffay@cs.unc.edu
October 4, 2017


<http://www.cs.unc.edu/~jeffay/courses/comp530>

©2017 by Kevin Jeffay

1

HW 4

A message passing producer/consumer system




- ◆ Re-implement your solution to HW 3 as a pipeline of Unix processes that communicate via message passing using Linux “pipes”
- ◆ Major changes to HW3:
 - » Each one of your ST threads will become a Linux process
 - » Each one of your bounded buffers will become a Linux pipe

©2017 by Kevin Jeffay

2

HW 4

Details — Linux Interprocess Communication (IPC)




- ◆ IPC in Linux is like message passing except more general
- ◆ The programming abstraction is a “pipe” — a shared, in-memory file
 - » A queue of 4K bytes
 - » Linux *read/write* systems calls used to “pass messages” between processes

©2017 by Kevin Jeffay

3

HW 4

Details — Linux Interprocess Communication (IPC)




- ◆ Pipes provide a form of buffered, asynchronous message passing
 - » Data written to one end of the pipe by one process can be read at the other end of the pipe by another process
 - ❖ A reader of the pipe blocks when the queue is empty
 - ❖ A writer to the pipe blocks when the queue is full
 - » Data is handled in a FIFO (first-in-first-out) order
 - » Pipes are unidirectional

©2017 by Kevin Jeffay

4

HW 4

Details — Linux Interprocess Communication (IPC)




- ◆ Typically, a process creates a pipe just before it forks one or more child processes
 - » Creating a pipe results in a process receiving a read file descriptor and a write file descriptor
- ◆ The pipe is then used for communication between the parent and child process, or between two sibling processes
 - » One process uses the read file descriptor (and ignores the write descriptor) and a second process uses the write file descriptor (and ignores the read descriptor)

©2017 by Kevin Jeffay

5

Linux Pipes

Programming details



- ◆ The Linux `pipe()` function:
 - » Is declared in the header file `unistd.h`
 - » `int pipe(int myPipe[2])`
 - ❖ Creates a pipe and puts the file descriptors for the reading and writing ends of the file (respectively) into `myPipe[0]` and `myPipe[1]`
 - ❖ If successful, pipe returns a value of 0
 - ❖ On failure, -1 is returned
 - » Functions `read()` and `write()` then used to send/receive messages

©2017 by Kevin Jeffay

6

Linux Pipes

Programming details




```
main() {
    int pipe1[2], pipe2[2];
    if (pipe(pipe1) == -1 || pipe(pipe2) == -1) error(...)
    switch (childpid = fork()) {
        case 0: /* child */
            close(pipe1[1]); /* write descriptor for pipe1 */
            close(pipe2[0]); /* read descriptor for pipe2 */
            client(pipe1[0], pipe2[1]); /* start client program */
        default: /* parent */
            close(pipe1[0]); /* read descriptor for pipe1 */
            close(pipe2[1]); /* write descriptor for pipe2 */
            server(pipe2[0], pipe1[1]); /* start server program */
            while (wait((int *) 0) != childpid); /* wait for child */
    }
}
```

©2017

7

Linux Pipes

More programming details



- ◆ You're going to be creating lots of processes in this assignment...
- ◆ Make sure to check for zombie processes you've created!
 - » `ps -ef | grep -e PID -e YOUR-LOGIN-NAME`
 - » `kill pid-number`
- ◆ Also, be sure to limit the maximum number of processes you can create

©2017 by Kevin Jeffay

8